

Thermodynamic Bounding for Cryptographic Prime Generation

The Adaptive Sigma Sieve

Drift Systems Inc.
driftsystems.io

December 2025 — TRL-5 FPGA-Verified — 63 Provisional Patent Applications Filed

Abstract

Cryptographic prime generation for 4096-bit keys is dominated by the computational cost of probabilistic primality tests (Miller–Rabin). We present the **Adaptive Sigma Sieve**, a pre-filter that rejects candidate integers based on their bit-structure metrics (Hamming weight, transition density) before they reach downstream modular exponentiation. By dynamically scaling acceptance bounds ($\mu \pm k\sigma$) with bit-length L , the filter shapes the candidate stream to reduce worst-case downstream cost. Empirical results on representative datasets demonstrate a 29–32% reduction in Miller–Rabin workload at typical operating points, increasing throughput per watt of HSM compute without changing the correctness authority (final primality testing remains Miller–Rabin).

1. Introduction

The generation of RSA and ECC primitives relies on “Blind Search”: generating random odd integers n and testing for primality. As key lengths L increase, the density of primes decreases ($1/\ln n$), while the cost of each Miller–Rabin round scales as $O(k \cdot L^3)$. Existing acceleration methods focus on modular sieving (wheel factorization) using small-prime divisibility. Wheel sieves ignore the *additive* structure of the candidate. The Adaptive Sigma Sieve is orthogonal to small-prime filtering: it operates on bit-level distribution metrics rather than divisibility, and can be composed with wheel sieves in series.

2. The Scaling Problem

For a uniform random integer of length L , the Hamming weight $H(n)$ follows a binomial distribution that approaches a normal distribution by Central Limit Theorem:

$$\mu = \frac{L}{2}, \quad \sigma = \frac{\sqrt{L}}{2} \quad (1)$$

Static percentage filters (e.g., accepting 45%–55% weight) fail at high bit depths ($L = 4096$) because the relative standard deviation shrinks with L , so a fixed-percentage window either rejects valid candidates or accepts most of the distribution. An adaptive bound that scales with σ is required.

3. Methodology: The Adaptive Sigma Sieve

We define a dynamic acceptance interval A_L centered on the binomial mean and scaled by the observed standard deviation:

$$A_L = \left[\frac{L}{2} - k \frac{\sqrt{L}}{2}, \frac{L}{2} + k \frac{\sqrt{L}}{2} \right] \quad (2)$$

where k is a tuning coefficient in $[0.8, 1.2]$ chosen to target a desired survivor rate (typically 10–50% depending on cost-per-prime trade-off).

3.1 Hardware Implementation

To compute Hamming weight at line rate for 4096-bit candidates, we use a **Split-Loop Adder Tree** on the Sipeed Tang Primer 20K (Gowin GW2A) FPGA. The architecture computes partial popcounts $S_{0..7}$ over disjoint candidate chunks in parallel and aggregates them in a single clock cycle, keeping synthesis depth bounded. The split-loop equivalence to a naive sum is formally verified in Lean 4 (`SplitLoopAdder.lean`; no sorry).

3.2 Transition Density (Orthogonal Metric)

Hamming weight alone does not detect structured bit-streams (e.g., long runs). We additionally compute transition density $T = \text{popcount}(N \oplus (N \gg 1))$ and

apply a parallel acceptance window. Transition density flags structured candidates invisible to weight alone, and helps normalize downstream modular-arithmetic switching behavior.

4. Results

In a sample of $N = 100,000$ candidates generated uniformly at random, the filter rejected candidates falling outside the configured acceptance windows:

- 512-bit candidates: **29.6%** reduction in Miller–Rabin workload at $k = 1.0$.
- 4096-bit candidates: **31.6%** reduction in Miller–Rabin workload at $k = 1.0$.

The gain is consistent across key sizes because both μ and σ scale with L , so the filter operates at the same relative position regardless of bit-depth.

5. Honest Framing

The structural gate is a throughput optimization layer, not a primality oracle:

- **What this paper shows:** Bit-structure pre-filtering reduces the average number of Miller–Rabin rounds required per prime, at FPGA line rate.
- **What this paper does *not* show:** A direct primality test, cryptographic-strength resistance against cryptanalysis, or asymptotic complexity improvement. Final acceptance testing (Miller–Rabin, or BPSW for additional confidence) remains the correctness authority.

6. Status

TRL-5 FPGA-verified on Tang Primer 20K (Gowin GW2A); cycle-accurate against the software model. Independent cryptanalytic peer review of the broader DAD core is pending.

Contact

Lukas Cain, Systems Architect
 licensing@driftsystems.io
<https://driftsystems.io>